

Optimizing Consumer Credit Risk Scoring: Leakage-Free Data Engineering, Calibrated Gradient Boosters, and Ensembled Probability of Default Estimation

Bhanu Prakash Bhavanam
Department of Computer
science engineering
KLH University
Hyderabad, India
2420030209@klh.edu.in

Krishita Myneni
Department of Computer
science engineering
KLH University
Hyderabad, India
2420030443@klh.edu.in

Laasya Priya Duggisetty
Department of Computer
science engineering
KLH University
Hyderabad, India
2420030002@klh.edu.in

Sunanda
Department of Computer
science engineering
KLH University
Hyderabad, India

Abstract—Accurate quantification of credit default risk is foundational to modern financial stability, capital allocation, and regulatory solvency under the Basel Accords. While contemporary machine learning literature frequently reports near-perfect classification performance on peer-to-peer lending benchmarks, such figures almost ubiquitously stem from subtle post-origination data leakage (e.g., training on post-issuance recoveries, payment trajectories, and debt settlement tags). In this paper, we propose a rigorous, leak-free, end-to-end machine learning credit risk modeling framework evaluated on 100,000 real-world consumer loans from the Lending Club portfolio. Our methodology integrates: (1) systematic auditing and removal of 40 post-origination leakage features; (2) a quantitative multi-criteria data quality framework improving portfolio health from 99.39% to 100.00%; (3) domain-driven financial interaction engineering capturing debt-burden velocity, delinquency severity, and inquiry pressure; (4) strict chronological out-of-time train/test partitioning (69,867 training vs. 17,467 testing loans); and (5) post-hoc Isotonic probability calibration applied to gradient boosted decision trees (XGBoost, LightGBM, CatBoost) and a Soft-Voting Champion Ensemble. Our tuned Champion models achieve an Area Under the Receiver Operating Characteristic (ROC-AUC) of 0.7488, a Kolmogorov-Smirnov (KS) statistic of 0.3640 (well surpassing the 0.30 banking deployment benchmark), an overall classification accuracy of 70.93%, a defaulter detection recall of 65.44%, and a top-decile lift of 3.21x (capturing 32.1% of all defaulters in the top 10% riskiest cohort). Finally, calibrated default probabilities are mapped to a standard 300–850 credit scorecard scale via Points-to-Double-the-Odds (PDO) formulation, bridging algorithmic precision with institutional interpretability and regulatory deployability.

Keywords—Credit Risk Scoring, Probability of Default (PD), Data Leakage, Gradient Boosted Decision Trees, XGBoost, LightGBM, CatBoost, Isotonic Calibration, Kolmogorov-Smirnov Statistic, Expected Loss, Scorecard Scaling.

I. INTRODUCTION

Credit risk—the probability that a borrower defaults on contractual debt obligations—represents the preeminent category of exposure faced by banking institutions, peer-to-

peer (P2P) lending platforms, and retail credit providers [1]. Under the Basel II and Basel III regulatory frameworks, credit risk assessment directly dictates minimum regulatory capital requirements through the Expected Loss (EL) equation:

$$EL = PD \times LGD \times EAD \quad (1)$$

where PD denotes the Probability of Default, LGD represents the Loss Given Default (the percentage of exposure unrecoverable post-charge-off), and EAD is the Exposure at Default [2]. Consequently, an estimation error of even 1–2 percentage points in PD compounds across multi-billion-dollar credit portfolios, precipitating severe balance-sheet impairment or mispriced credit terms [3].

Historically, retail lenders relied on statistical scorecards underpinned by multivariable Logistic Regression and Discriminant Analysis, prized for their parametric transparency, monotonicity, and ease of audit [4]. However, traditional linear approaches inherently fail to capture complex non-linear interactions, credit-line saturation cliffs, and asymmetric default triggers characteristic of consumer financial behavior. While modern non-parametric algorithms—specifically tree-based ensembles—offer superior functional flexibility, academic and industrial implementations are frequently compromised by two systemic pitfalls:

1) **Pervasive Data Leakage:** Over 60% of open-source studies on peer-to-peer datasets mistakenly incorporate features generated after loan issuance (such as total payments received, last payment dates, and debt collection recovery fees), yielding artificially inflated ROC-AUC values exceeding 0.95 that collapse in production underwriting [5].

2) **Probability Distortion via Class Imbalance Correction:** Because defaulters constitute a minority class (~15–20% of settled portfolios), practitioners routinely introduce synthetic sampling (e.g., SMOTE) or class weighting (scale_pos_weight). While this sharpens rank-order discrimination, it severely warps the class prior, producing uncalibrated probabilities where a borrower with a true 8%

default propensity is scored at 35%, completely invalidating regulatory capital and pricing models [6].

To resolve these challenges, this paper presents an end-to-end, scientifically defensible credit scoring framework. The primary contributions of this work are threefold:

- **Leakage-Free Data Sanitization & Audit:** We establish an objective exclusion protocol eliminating 40 post-origination variables, reinforced by quantitative multi-criteria data quality scoring before and after automated remediation.
- **Domain-Specific Financial Feature Engineering:** We construct novel interaction features reflecting consumer debt velocity, amortization strain, and credit-bureau inquiry velocity.
- **Calibrated Multi-Model Benchmarking:** We train, tune, and compare Logistic Regression, Random Forest, XGBoost, LightGBM, CatBoost, and a Soft-Voting Champion Ensemble over a chronological (out-of-time) split, applying post-hoc Isotonic Regression to guarantee well-calibrated, production-ready default probabilities.
- **Regulatory Scorecard & Policy Banding:** We implement a Points-to-Double-the-Odds (PDO) mathematical transformation to convert calibrated default probabilities into standard 300–850 credit scores with tiered underwriting protocols.

II. RELATED WORK & THE DATA LEAKAGE TRAP

A. Evolution of Credit Scoring Methodology

Credit scoring methodologies originated with Fisher's Linear Discriminant Analysis and Altman's Z-score formulation for corporate solvency [7]. In consumer lending, Thomas et al. [1] and Hand and Henley [4] formalized logistic regression scorecards using Weight of Evidence (WoE) transformations. Subsequent literature introduced tree ensembles; Kruppa et al. demonstrated that Random Forests reduce classification errors relative to scorecards [8]. More recently, Gradient Boosted Decision Trees (GBDTs)—specifically XGBoost [9], LightGBM [10], and CatBoost [11]—have emerged as state-of-the-art benchmarks in tabular financial domains due to their capacity to handle high-dimensional collinearity and sparse categorical indicators.

B. The Post-Origin Data Leakage Phenomenon

A pervasive deficiency across peer-to-peer credit risk literature (particularly studies utilizing the Lending Club repository) is temporal data leakage. Lending Club datasets ship with over 150 attributes. However, roughly 40 of these variables (e.g., `total_pymnt`, `recoveries`, `last_pymnt_d`, `collection_recovery_fee`, `out_prncp`) are updated throughout the repayment lifespan of the loan. If an algorithm is trained on recoveries (which is non-zero almost exclusively when a loan enters charge-off status), the model trivializes the learning problem by retroactively inferring the outcome rather

than forecasting future default propensity from applicant traits known strictly at origination ($t = 0$). Published studies reporting ROC-AUC figures of 0.92–0.99 almost invariably harbor leakage. As documented by industry practitioners and regulatory examiners, a leak-free, deployable consumer credit model operating solely on origination data realistically achieves ROC-AUC scores between 0.68 and 0.75 [3], [5]. Models exceeding 0.85 on raw borrower application data warrant immediate audit for backward information transfer.

C. Probability Calibration in Credit Risk

Supervised classifiers trained on re-weighted or subsampled credit datasets output uncalibrated risk scores that cannot be interpreted as posterior probabilities $P(Y=1|x)$ [6]. In banking, uncalibrated scores disrupt Expected Loss calculations and capital reserves. Calibration methods, notably Platt Scaling (logistic) and Isotonic Regression (non-parametric monotonic step functions), map raw scores to true empirical event frequencies [15]. In this work, we demonstrate that post-hoc Isotonic Regression strictly preserves ranking power while driving the Brier score down to optimal levels.

III. DATA ENGINEERING & SANITIZATION PROTOCOL

A. Data Collection and Integration

Experiments were conducted using the comprehensive Lending Club loan database spanning 2007 through 2018Q4. To maintain computational tractability while preserving asymptotic statistical power, an integrated cohort of 100,000 real loans was loaded across 42 origination attributes.

B. Audit of Post-Origin Data Leakage Features

An exhaustive institutional audit identified and purged 40 post-origination columns prior to model training. Key excluded categories include:

- **Repayment History:** `total_pymnt`, `total_pymnt_inv`, `total_rec_prncp`, `total_rec_int`, `total_rec_late_fee`.
- **Collection and Charge-off Tags:** `recoveries`, `collection_recovery_fee`, `last_pymnt_d`, `last_pymnt_amnt`, `next_pymnt_d`.
- **Outstanding Balances:** `out_prncp`, `out_prncp_inv`.
- **Post-Issuance Bureau Updates:** `last_credit_pull_d`, `hardship_flag`, `settlement_status`, `debt_settlement_flag`.

C. Data Quality Audit Framework

Data hygiene was quantitatively evaluated before and after automated cleaning across four core quality dimensions, defined with institutional weighting:

$$Q_{\text{overall}} = w_c Q_c + w_v Q_v + w_s Q_s + w_u Q_u \quad (2)$$

where $w_c = 0.35$ (completeness), $w_v = 0.25$ (validity), $w_s = 0.20$ (consistency), and $w_u = 0.20$ (uniqueness).

As summarized in Table I, raw data exhibited missingness in delinquency timelines (mths_since_last_delinq at 48.19%), employment spans (emp_length at 6.11%), and bankcard utilization (bc_util at 1.03%), alongside 644 near-duplicate rows. Post-cleaning remediation elevated overall quality to 100.00% (Grade A+).

TABLE I. DATA QUALITY AUDIT BEFORE AND AFTER REMEDIATION

Quality Dimension	Before Clean	After Clean	Delta
Completeness	98.63%	100.00%	+1.37%
Validity	100.00%	100.00%	0.00%
Consistency	100.00%	100.00%	0.00%
Uniqueness	99.36%	100.00%	+0.64%
OVERALL SCORE	99.39 / 100	100.00 / 100	+0.61
Grade	A (Excellent)	A+ (Perfect)	Remediated

D. Cleaning and Sanitization Steps

1) Target Definition & Cohort Resolution: Raw loan_status was mapped to Default (Y=1) including Charged Off and Default, and Non-Default (Y=0) including Fully Paid. Indeterminate loans (Current, Late 16-120 days, In Grace Period) were excised (12,108 loans dropped), leaving 87,892 fully resolved loans with an empirical portfolio default rate of 20.03%.

2) String Parsing & Deduplication: Term was parsed from ' 36 months' to integer 36; interest rates and revol_util were stripped of percentage signs; emp_length was mapped to integers 0-10. 558 near-duplicate records sharing identical applicant demographic vectors were pruned, yielding 87,334 clean unique borrowers.

3) Outlier Winsorization: Continuous debt ratios and loan amounts were winsorized at 0.5% and 99.5% quantiles, followed by median imputation calculated strictly over historical training partitions.

IV. FEATURE ENGINEERING & CHRONOLOGICAL PARTITIONING

A. Engineered Financial Indicators

To capture non-linear risk signals beyond linear application fields, 17 interaction features were derived:

1) Affordability & Leverage: Loan-to-Income (LTI = loan_amnt / annual_inc), Payment Burden (installment / monthly_inc), and Interest Rate Pressure (dti × int_rate).

2) Bureau Velocity: High Revolving Utilization Flag (revol_util > 80%), Inquiry Velocity (inq_last_6mths /

credit_years), and Active Account Share (open_acc / total_acc).

3) Insolvency Severity Index: Weighted composite metric synthesizing multi-tiered credit impairment:

$$Severity = 2.0 \text{ delinq} + 1.5 \text{ pub_rec} + 3.0 \text{ bankruptcies} \quad (3)$$

4) Temporal Bureau Depth: Credit history tenure computed from issue date to earliest credit line in months.

Following expansion, the modeling space comprised 31 numeric predictors and 8 categorical variables (59 one-hot encoded columns).

B. Chronological (Out-of-Time) Validation Split

Standard k-fold cross-validation introduces future lookahead bias by shuffling across credit cycles. We enforce a strict chronological partition: loans issued prior to December 2015 serve as the training partition (N = 69,867), while loans issued from December 2015 onward serve as the out-of-time test cohort (N = 17,467).

V. MACHINE LEARNING ARCHITECTURE & CALIBRATION

A. Model Zoo Specifications

We benchmarked six standardized scikit-learn pipeline architectures: (1) Logistic Regression with L2 regularization; (2) Random Forest (250 trees, max_depth=12); (3) XGBoost (n_estimators=300, max_depth=5, lr=0.04, colsample=0.8); (4) LightGBM (num_leaves=31, max_depth=6, lr=0.035); (5) CatBoost (iterations=300, depth=5, lr=0.05); and (6) a Champion Soft-Voting Ensemble aggregating predictions across XGBoost, LightGBM, and Random Forest:

$$P_{Ensemble}(Y=1|x) = 1/3 [P_{XGB}(x) + P_{LGB}(x) + P_{RF}(x)] \quad (4)$$

B. Class Imbalance Mitigation

To avoid the spatial distortion of SMOTE on high-dimensional financial data, class imbalance (4:1 non-default to default) was addressed algorithmically via loss weighting:

$$scale_pos_weight = N_{neg} / N_{pos} \approx 3.99 \quad (5)$$

C. Post-Hoc Isotonic Probability Calibration

While class reweighting sharpens ROC-AUC rank-ordering, it systematically inflates raw predicted probabilities above empirical priors. We partitioned the training set into an 85% fitting partition (N = 59,386) and a 15% calibration partition (N = 10,481). Post-hoc Isotonic Regression minimizes squared errors subject to monotonicity constraints. Crucially, because Isotonic Regression is strictly monotonic, ROC-AUC and KS statistics are mathematically invariant, while Brier score drops from 0.179 to 0.1454.

D. Decision Threshold Optimization

The standard default threshold of 0.50 collapses recall in imbalanced credit portfolios. We swept thresholds to

maximize default F1-score, converging to optimal operating cut-offs at τ^* in [0.19, 0.23].

VI. EXPERIMENTAL RESULTS & BENCHMARK EVALUATION

Models were evaluated on the holdout out-of-time test partition (N = 17,467) across ROC-AUC, Kolmogorov-Smirnov (KS), PR-AUC, Accuracy, Recall, Precision, and Brier score.

TABLE II. COMPREHENSIVE MULTI-MODEL BENCHMARK RESULTS (N = 17,467)

Model	ROC-AUC	KS Stat	PR-AUC	Accuracy	Recall	F1 (Def)	Brier
XGBoost (Tuned)	0.7488	0.3631	0.4457	70.93%	62.97%	0.4830	0.1455
Champion Ensemble	0.7487	0.3640	0.4458	69.77%	65.44%	0.4829	0.1454
CatBoost	0.7481	0.3614	0.4446	70.05%	64.59%	0.4819	0.1456
LightGBM	0.7476	0.3652	0.4462	68.44%	67.75%	0.4808	0.1458
Logistic Regression	0.7448	0.3591	0.4385	68.62%	66.79%	0.4786	0.1463
Random Forest	0.7443	0.3568	0.4385	70.73%	62.33%	0.4788	0.1462




Fig. 1. Out-of-time Receiver Operating Characteristic (ROC) curves across all benchmark models on holdout test data.




Fig. 2. Reliability calibration curves before and after post-hoc Isotonic Regression calibration.

As documented in Table II and Fig. 1, tuned XGBoost and the Champion Ensemble demonstrated superior discrimination, achieving ROC-AUC of 0.7488 and KS of 0.3640, outperforming the baseline Logistic Regression (0.7448) and Random Forest (0.7443). Furthermore, at $\tau^* = 0.21$, the Champion Ensemble caught 65.44% of all defaulting borrowers.

TABLE III. DECILE LIFT & CUMULATIVE GAINS VALIDATION

Decile	Loans	Defaults	Actual Rate	Avg PD	Lift	Cum Def %
1 (High)	1,747	842	48.20%	35.8%	3.21x	32.1%
2	1,747	571	32.68%	25.1%	2.18x	53.9%
3	1,746	445	25.49%	20.4%	1.70x	70.9%
4	1,747	332	19.00%	16.8%	1.27x	83.6%
5	1,747	218	12.48%	13.7%	0.83x	91.9%
6	1,747	124	7.10%	10.9%	0.47x	96.6%
7	1,747	58	3.32%	8.3%	0.22x	98.8%
8	1,746	20	1.15%	6.2%	0.08x	99.6%
9	1,747	8	0.46%	4.5%	0.03x	99.9%
10 (Low)	1,747	2	0.11%	2.8%	0.01x	100.0%
Portfolio	17,467	2,620	15.00%	14.5%	1.00x	100.0%



Fig. 3. Relative feature importance rankings across top credit default risk drivers.




Fig. 4. Decile Lift and Cumulative Default capture rates showing 3.21x top-decile concentration.

C. Decile Lift and Business Validation

As detailed in Table III and Fig. 4, sorting test applicants into decile risk buckets yielded exceptional separation: Decile 1 had an actual default rate of 48.20% (Lift = 3.21x vs. the 15.00% book average). The top two deciles captured 53.9% of all portfolio defaults, while Decile 10 suffered merely 0.11% defaults. Furthermore, feature importance (Fig. 3) showed `dti_x_int_rate` (18.4%), `int_rate` (14.2%), `installment_to_monthly_income` (11.6%), and `fico_avg` (9.8%) as primary risk drivers.

VII. CREDIT SCORECARD TRANSFORMATION & RISK BANDING

A. Probability to FICO-Score Mapping

We employ the industry standard Points-to-Double-the-Odds (PDO) credit scorecard transformation:

$$\text{Odds} = (1 - \text{PD}) / \text{PD} \quad (6)$$

$$\text{Score} = \text{Offset} + \text{Factor} \times \ln(\text{Odds}) \quad (7)$$

where $\text{Factor} = \text{PDO} / \ln(2)$ and $\text{Offset} = \text{Base Score} - \text{Factor} \times \ln(\text{Base Odds})$. Setting Base Score = 660, PDO = 40, and Base Odds = 4.0 (corresponding to 20% PD), scores are mapped to the interval [300, 850].

TABLE IV. INSTITUTIONAL RISK BAND POLICY ALLOCATION

Risk Tier	Default Prob (PD)	Score Range	Portfolio %	Observed PD	Underwriting Action
-----------	-------------------	-------------	-------------	-------------	---------------------

Low Risk	PD ≤ 10.0%	700 - 850	35.0%	6.4%	Auto-Approve (Prime)
Medium Risk	10.0% < PD ≤ 25.0%	640 - 699	47.0%	16.8%	Manual Review / Limits
High Risk	PD > 25.0%	300 - 639	18.0%	36.2%	Decline / Collateral

VIII. CONCLUSION

In this study, we formulated, implemented, and validated an end-to-end, leakage-free credit risk scoring pipeline on real-world Lending Club loan data. By purging 40 post-origination leakage attributes, engineering 17 non-linear debt velocity features, enforcing chronological out-of-time validation, and pairing gradient boosted decision trees with post-hoc Isotonic Probability Calibration, our models achieve high discrimination (ROC-AUC 0.7488, KS 0.3640) and strong operational efficiency (70.93% Accuracy, 65.44% Default Recall, 3.21x Top-Decile Lift). Furthermore, integrating Points-to-Double-the-Odds scorecard transformation ensures direct compatibility with institutional risk management and Basel capital calculations.

REFERENCES

- [1] Thomas, Lyn & Edelman, David & Crook, Jonathan. (2002). Credit Scoring and its Applications.
- [2] Basel Committee on Banking Supervision. 2006. International Convergence of Capital Measurement and Capital Standards: A Revised Framework (Comprehensive Version). Basel: Bank for International Settlements, June. bis.org
- [3] E. I. Altman and A. Saunders, 'Credit risk measurement: Developments over the last 20 years,' Journal of Banking & Finance, vol. 21, pp. 1721–1742, 1997.
- [4] D. J. Hand and W. E. Henley, 'Statistical classification methods in consumer credit scoring: A review,' JRSS Series A, vol. 160, pp. 523–541, 1997.
- [5] S. Lessmann, B. Baesens, H. Seow, and L. C. Thomas, 'Benchmarking classification algorithms for credit scoring,' EJOR, vol. 247, pp. 124–136, 2015.
- [6] A. Niculescu-Mizil and R. Caruana, 'Predicting good probabilities with supervised learning,' in Proc. ICML, Bonn, Germany, 2005, pp. 625–632.
- [7] E. I. Altman, 'Financial ratios, discriminant analysis and corporate bankruptcy,' The Journal of Finance, vol. 23, pp. 589–609, 1968.
- [8] J. Kruppa et al., 'Consumer credit risk: Probability of default using classification trees and random forests,' Acta Informatica Medica, vol. 21, 2013.
- [9] T. Chen and C. Guestrin, 'XGBoost: A scalable tree boosting system,' in Proc. KDD, San Francisco, CA, 2016, pp. 785–794.
- [10] G. Ke et al., 'LightGBM: A highly efficient gradient boosting decision tree,' in NeurIPS, Long Beach, CA, 2017, pp. 3146–3154.
- [11] L. Prokhorenkova et al., 'CatBoost: unbiased boosting with categorical features,' in NeurIPS, Montreal, Canada, 2018, pp. 6638–6648.
- [12] J. Huang, J. Lu, and C. X. Ling, 'Comparing naive Bayes, decision trees, and SVM with AUC and accuracy,' in Proc. IEEE ICDM, 2003, pp. 553–556.
- [13] F. Provost and T. Fawcett, 'Analysis and visualization of classifier performance under imprecise class distributions,' in Proc. KDD, 1997.
- [14] M. Schreiner, 'Credit scoring for microfinance: Is it worth it?,' Journal of Microfinance, vol. 4, no. 1, pp. 105–124, 2002.
- [15] B. Zadrozny and C. Elkan, 'Transforming classifier scores into accurate multiclass probability estimates,' in Proc. KDD, 2002, pp. 694–699.